

Anvesh Sai Vemuri

Jersey City, NJ | +1 315-420-9623 | anveshsvemuri@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data Engineer with 4+ years of experience designing and building scalable ETL/ELT pipelines, distributed data processing systems, data lakes, lakehouse architectures, and cloud data warehouses using Python, SQL, PySpark, Apache Spark, AWS, and Databricks. Experienced in data modeling, API-driven ingestion, workflow orchestration, data quality, and performance optimization across large-scale enterprise datasets. Hands-on experience with Amazon S3, AWS Glue, Lambda, Redshift, EMR, Snowflake, Airflow, dbt, and Kafka.

SKILLS

Languages: Python, SQL, Java

Data Engineering: Apache Spark, PySpark, Databricks, Apache Airflow, dbt, Pandas, NumPy, ETL/ELT, REST APIs, Delta Lake, Dimensional Data Modeling, Kafka, Hadoop, Hive, HDFS, Boto3

Cloud & Databases: AWS (Amazon S3, AWS Glue, AWS Lambda, Amazon Redshift, Athena, AWS Lake Formation, Amazon EMR), Snowflake, MySQL, PostgreSQL

Tools & Visualization: Git, GitHub, Azure DevOps, Jenkins, Docker, Linux, Alteryx, Jupyter Notebook, Tableau, Power BI

EXPERIENCE

Publicis Groupe | *Data Engineer*

Sep 2024 - Aug 2026

- Designed and maintained scalable ETL/ELT pipelines using Python, SQL, Databricks, Amazon S3, Amazon Redshift, and REST APIs to ingest and transform advertising and eCommerce data from 20+ external platforms, reducing reporting latency by 40% while supporting analytics for 6+ enterprise clients.
- Reverse-engineered and migrated 50+ legacy Alteryx workflows to Databricks, rebuilding transformation logic in PySpark and SQL and reducing execution time by 35% while improving scalability, maintainability, and deployment efficiency.
- Developed and optimized dbt and SQL transformations and Amazon Redshift stored procedures for incremental processing, large-scale aggregations, and analytics reporting, improving query performance by 30% while increasing warehouse efficiency and pipeline reliability.
- Developed automated data quality, monitoring, and alerting frameworks using Python and SQL to detect pipeline failures, schema changes, reconciliation issues, and source-to-target discrepancies, reducing manual QA effort by 10 hours per week. Used Azure DevOps/Git and CI/CD workflows for version control, peer code reviews, testing, and controlled deployment of production data pipeline changes.
- Developed reusable Python ingestion pipelines integrating data from 20+ third-party platforms through REST APIs, SFTP, and cloud storage into Databricks and Amazon Redshift, supporting standardized analytics and reporting across 100+ dashboards.

JPMorgan Chase & Co. | *Data Engineer*

Sep 2023 - Aug 2024

- Developed distributed PySpark ETL pipelines processing 20M+ financial and transactional records daily, improving processing performance by 45% while supporting downstream analytics and reporting.
- Built 100+ SQL and dbt transformations in Snowflake and Hive for large-scale financial and transactional datasets, implementing reusable data models and improving analytical query performance by 35%.
- Optimized Spark and Hive ETL workloads by implementing partition-aware processing, leveraging Parquet storage formats, and tuning Spark SQL transformations to improve performance and scalability for large enterprise datasets.
- Automated and maintained 25+ Apache Airflow workflows for enterprise data pipelines, improving orchestration, monitoring, and data reliability while reducing manual intervention by 12 hours per week.
- Implemented automated dbt tests and data quality checks to validate schema integrity, transformation logic, and downstream reporting accuracy across financial datasets.

Dixon Technologies | *Data Engineer*

Aug 2020 - Jul 2021

- Developed 15+ custom Apache Airflow operators in Python integrating Snowflake, Slack, and Tableau workflows, automating enterprise data pipelines and improving monitoring and operational efficiency.
- Designed and implemented Amazon Redshift data warehouse schemas and ETL pipelines using Python and SQL, improving query performance by 35% and reducing report generation time by 40% for high-volume analytics workloads.
- Built event-driven Spark Streaming and AWS Lambda pipelines processing 5M+ real-time events daily, enabling low-latency anomaly detection and improving operational responsiveness.
- Developed scalable ETL pipelines using AWS Glue, Amazon S3, AWS Lambda, Boto3, Databricks, and PySpark, leveraging Boto3 for automated S3 file ingestion and object management, reducing cloud compute costs by 25% while improving pipeline efficiency and scalability.

EDUCATION

Southwest Baptist University | *PhD in Technology and Artificial Intelligence*

Sep 2026 - Present

New England College | *Master of Science, Computer Information Systems*

Dec 2023

CERTIFICATIONS

- HackerRank SQL (Advanced) Certificate | HackerRank | 2026
- Databricks Academy – Advanced Data Engineering Training (Unity Catalog, Lakeflow Jobs, Lakeflow Connect, Spark Declarative Pipelines) | 2025