

Anvesh Sai Vemuri

Jersey City, NJ | +1 315-420-9623 | anveshsvemuri@gmail.com | [LinkedIn](#) | [GitHub](#)

SUMMARY

Data Engineer with 4+ years of experience designing scalable ETL/ELT pipelines, distributed data processing systems, and cloud-native analytics platforms using Python, SQL, PySpark, Apache Spark, AWS, and Databricks. Experienced in building modern data lakes, data warehouses, scalable analytics solutions, and reliable data models for enterprise analytics. Hands-on experience with Amazon S3, AWS Glue, AWS Lambda, Amazon Redshift, Snowflake, and modern data integration technologies, with a strong focus on automation, API-driven data pipelines, data quality, performance optimization, and AI-powered data engineering solutions.

SKILLS

Languages: Python, SQL, Java

Data Engineering: Apache Spark, PySpark, Databricks, Apache Airflow, dbt, Pandas, NumPy, ETL/ELT, REST APIs, Data Modeling, Delta Lake, Kafka, Hadoop, Hive, HDFS, Boto3

Cloud & Databases: AWS (Amazon S3, AWS Glue, AWS Lambda, Amazon Redshift, Athena, AWS Lake Formation), Snowflake, MySQL, PostgreSQL

DevOps & Tools: GitHub, Git, Jupyter Notebook, Jenkins, Docker, Linux

Visualization: Power BI, Tableau

EXPERIENCE

Publicis Groupe | Data Engineer

Sep 2024 - Present

Clients: Lactalis | Reckitt | J.M. Smucker | GST

- Designed and maintained scalable ETL/ELT pipelines using Python, SQL, Databricks, Amazon S3, Amazon Redshift, and REST APIs to ingest, transform, and operationalize data from 20+ external platforms, reducing reporting latency by 40% while supporting analytics for 6+ enterprise clients.
- Led migration of **50+** legacy Alteryx workflows to Databricks by redesigning transformation logic in PySpark and SQL, reducing execution time by **35%** while improving scalability, deployment efficiency, and long-term maintainability.
- Developed reusable Python-based ingestion and ETL pipelines integrating data from 20+ third-party platforms through APIs, SFTP, and cloud storage into Databricks and Amazon Redshift, enabling standardized reporting across 100+ dashboards while reducing manual data preparation by 50%.
- Engineered complex SQL transformations and Amazon Redshift stored procedures for campaign reporting, incremental data processing, and large-scale aggregations, improving query performance by 30% while increasing production pipeline scalability and reliability.
- Developed automated data quality, reconciliation, and monitoring frameworks using Python and SQL to validate source-to-target transformations, detect schema changes, identify mapping issues, and reduce manual QA effort by 10 hours per week while improving production data reliability.

JPMorgan Chase & Co. | Data Engineer

Sep 2023 - Aug 2024

- Developed distributed PySpark ETL pipelines processing over 20 million enterprise records daily from transactional systems, improving processing performance by 45% while supporting downstream analytics and reporting.
- Built **100+** HiveQL and Snowflake SQL transformations supporting enterprise-scale financial datasets, improving analytical query performance by **35%**.
- Designed dimensional data models supporting **30+** downstream reporting datasets and enterprise analytics applications.
- Developed scalable Spark SQL and dbt transformation pipelines powering **25+** production workflows and enterprise analytics applications.
- Automated **25+** Airflow workflows integrated with AWS Lake Formation, reducing manual intervention by **12 hours/week** while strengthening data governance.

Dixon Technologies | Data Engineer

Aug 2020 - Jul 2021

- Developed **15+** custom Apache Airflow operators in Python to integrate Snowflake, Slack, and Tableau workflows, automating enterprise data pipelines and reducing manual integration effort by **50%**.
- Designed and implemented Amazon Redshift data warehouse schemas and ETL pipelines using Python and SQL, improving query performance by **35%** and reducing report generation time by **40%** for high-volume analytics workloads.
- Built Spark Streaming and AWS Lambda pipelines processing **5M+** real-time events daily, enabling low-latency anomaly detection and improving operational responsiveness.
- Developed scalable ETL pipelines using AWS Glue, Amazon S3, AWS Lambda, Boto3, Databricks, and PySpark, leveraging Boto3 for automated S3 file ingestion and object management, reducing cloud compute costs by 25% while improving pipeline efficiency and scalability.

EDUCATION

New England College

Master of Science, Computer Information Systems

Dec 2023

Henniker, NH

CERTIFICATIONS

- HackerRank SQL (Advanced) Certificate | HackerRank | 2026
- Databricks Academy – Advanced Data Engineering Training (Unity Catalog, Lakeflow Jobs, Lakeflow Connect, Spark Declarative Pipelines) | 2025