

Anvesh Sai Vemuri

Jersey City, NJ | +1 3154209623 | anveshsvemuri@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data Engineer with 4+ years of experience designing and building scalable ETL/ELT pipelines, distributed data processing systems, data lakes, lakehouse architectures, and cloud data warehouses using Python, SQL, PySpark, Apache Spark, AWS, and Databricks. Experienced in data modeling, API-driven ingestion, workflow orchestration, data quality, and performance optimization across large-scale enterprise datasets. Hands-on experience with Amazon S3, AWS Glue, Lambda, Redshift, EMR, Snowflake, Airflow, dbt, and Kafka.

SKILLS

- **Languages:** Python, SQL, Java
- **Data Engineering:** Apache Spark, PySpark, Databricks, Apache Airflow, dbt, Pandas, NumPy, ETL/ELT, REST APIs, Delta Lake, Dimensional Data Modeling, Kafka, Hadoop, Hive, HDFS, Boto3
- **Cloud & Databases:** AWS (Amazon S3, AWS Glue, AWS Lambda, Amazon Redshift, Athena, AWS Lake Formation, Amazon EMR), GCP, Snowflake, MySQL, PostgreSQL, BigQuery
- **Tools & Visualization:** Git, GitHub, Azure DevOps, Jenkins, Docker, Linux, Alteryx, Jupyter Notebook, Tableau, Power BI

EXPERIENCE

Publicis Groupe (Zenith Media Services) | *Data Engineer*

Sep 2024 - Aug 2026

- Engineered Python and SQL data pipelines ingesting and transforming advertising and eCommerce data from **20+ external platforms** across **BigQuery, Redshift, and Databricks**, reducing reporting latency by **40%** and supporting analytics for **6+ enterprise clients**.
- Migrated **50+ production Alteryx workflows to Databricks**, redesigning transformations and orchestration for scalable cloud processing and reducing execution time by **35%**.
- Optimized complex SQL transformations, stored procedures, and incremental data-processing workflows across cloud data warehouses, improving query performance by **30%** while supporting high-volume analytics workloads.
- Implemented automated data-quality checks, reconciliation logic, schema validation, and pipeline monitoring across production ETL workflows, reducing manual QA effort by **10+ hours per week** and improving reliability of downstream reporting.
- Integrated data from **20+ advertising and commerce platforms**, including **Pinterest, DoubleVerify, DV360/YouTube, Amazon, Walmart, Meta, TikTok, Snapchat, Reddit, and TTD**, standardizing campaign and performance data used across **100+ downstream dashboards**.

JPMorgan Chase & Co. | *Data Engineer*

Sep 2023 - Aug 2024

- Developed distributed PySpark ETL pipelines processing 20M+ financial and transactional records daily, improving processing performance by 45% while supporting downstream analytics and reporting.
- Built 100+ SQL and dbt transformations in Snowflake and Hive for large-scale financial and transactional datasets, implementing reusable data models and improving analytical query performance by 35%.
- Optimized Spark and Hive ETL workloads by implementing partition-aware processing, leveraging Parquet storage formats, and tuning Spark SQL transformations to improve performance and scalability for large enterprise datasets.
- Automated and maintained 25+ Apache Airflow workflows for enterprise data pipelines, improving orchestration, monitoring, and data reliability while reducing manual intervention by 12 hours per week.
- Implemented automated dbt tests and data quality checks to validate schema integrity, transformation logic, and downstream reporting accuracy across financial datasets.

Dixon Technologies | *Data Engineer*

Aug 2020 - Jul 2021

- Developed 15+ custom Apache Airflow operators in Python integrating Snowflake, Slack, and Tableau workflows, automating enterprise data pipelines and improving monitoring and operational efficiency.
- Designed and implemented Amazon Redshift data warehouse schemas and ETL pipelines using Python and SQL, improving query performance by 35% and reducing report generation time by 40% for high-volume analytics workloads.
- Built event-driven Spark Streaming and AWS Lambda pipelines processing 5M+ real-time events daily, enabling low-latency anomaly detection and improving operational responsiveness.
- Developed scalable ETL pipelines using AWS Glue, Amazon S3, AWS Lambda, Boto3, Databricks, and PySpark, leveraging Boto3 for automated S3 file ingestion and object management, reducing cloud compute costs by 25% while improving pipeline efficiency and scalability.

EDUCATION

Southwest Baptist University | *PhD, Technology and Artificial Intelligence*

Sep 2026 - Present

New England College | *Master of Science, Computer Information Systems*

Dec 2023

CERTIFICATIONS

- HackerRank SQL (Advanced) Certificate: HackerRank | 2026
- Databricks Academy – Advanced Data Engineering Training (Unity Catalog, Lakeflow Jobs, Lakeflow Connect, Spark Declarative Pipelines): 2025